

AI Featuring Engineering Assignment

Name: Krishna Kushwah

PRN: 20240110039

Batch: A2(AIML)

Dataset: IBM HR Analytics Attrition Dataset

```
# -----
# STEP 1: Install Kaggle API (only needed once per runtime)
# -----
!pip install -q kaggle
import os

# Directly store Kaggle API key (avoid manual upload every time)
os.environ['KAGGLE_USERNAME'] = "krishnakus" # ♦ Replace with your Kaggle username
os.environ['KAGGLE_KEY'] = "1c17326f6c19a146780799da2336ac06" # ♦ Replace with your actual Kaggle API key

# -----
# STEP 2: Download IBM HR Analytics dataset from Kaggle
# -----
!kaggle datasets download -d pavansubhasht/ibm-hr-analytics-attrition-dataset

# -----
# STEP 3: Unzip dataset file
# -----
!unzip -o ibm-hr-analytics-attrition-dataset.zip
```

```
Dataset URL: https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset
License(s): DbCL-1.0
ibm-hr-analytics-attrition-dataset.zip: Skipping, found more recently modified local copy (use --force to force download)
Archive: ibm-hr-analytics-attrition-dataset.zip
inflating: WA_Fn-UseC_-HR-Employee-Attrition.csv
```

```
# -----
# Import all required libraries
# -----
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
import warnings
warnings.filterwarnings('ignore')
```

```
# -----
# Load the dataset into a DataFrame
# -----
df = pd.read_csv('/content/WA_Fn-UseC_-HR-Employee-Attrition.csv')

# Display first 5 rows
df.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	Empl
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

```

# -----
# Check general information about the dataset
# -----
print("◆ Basic Info:")
df.info()

print("\n◆ Dataset Shape:", df.shape)
print("\n◆ Missing Values per Column:")
print(df.isnull().sum())

print("\n◆ Duplicate Rows:", df.duplicated().sum())


◆ Basic Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object  
 2   BusinessTravel   1470 non-null    object  
 3   DailyRate        1470 non-null    int64  
 4   Department       1470 non-null    object  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education        1470 non-null    int64  
 7   EducationField   1470 non-null    object  
 8   EmployeeCount    1470 non-null    int64  
 9   EmployeeNumber   1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object  
 12  HourlyRate       1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object  
 16  JobSatisfaction  1470 non-null    int64  
 17  MaritalStatus     1470 non-null    object  
 18  MonthlyIncome    1470 non-null    int64  
 19  MonthlyRate      1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    object  
 22  Overtime          1470 non-null    object  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours    1470 non-null    int64  
 27  StockOptionLevel  1470 non-null    int64  
 28  TotalWorkingYears 1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance  1470 non-null    int64  
 31  YearsAtCompany   1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64  
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

◆ Dataset Shape: (1470, 35)

◆ Missing Values per Column:
Age                0
Attrition          0
BusinessTravel     0
DailyRate           0
Department          0
DistanceFromHome   0
Education           0
EducationField      0
EmployeeCount       0
EmployeeNumber      0
EnvironmentSatisfaction 0

```

The dataset contains employee details for attrition prediction. It has no missing values or duplicate rows. Data types are mostly numeric, with some categorical fields like "Gender", "JobRole", and "MaritalStatus".

```

# -----
# Statistical Summary of all columns
# -----
df.describe(include='all').T

```

	count	unique	top	freq	mean	std	min	25%	50%	75%	max
Age	1470.0	NaN	NaN	NaN	36.92381	9.135373	18.0	30.0	36.0	43.0	60
Attrition	1470	2	No	1233	NaN	NaN	NaN	NaN	NaN	NaN	Na
BusinessTravel	1470	3	Travel_Rarely	1043	NaN	NaN	NaN	NaN	NaN	NaN	Na
DailyRate	1470.0	NaN	NaN	NaN	802.485714	403.5091	102.0	465.0	802.0	1157.0	1499
Department	1470	3	Research & Development	961	NaN	NaN	NaN	NaN	NaN	NaN	Na
DistanceFromHome	1470.0	NaN	NaN	NaN	9.192517	8.106864	1.0	2.0	7.0	14.0	29
Education	1470.0	NaN	NaN	NaN	2.912925	1.024165	1.0	2.0	3.0	4.0	5
EducationField	1470	6	Life Sciences	606	NaN	NaN	NaN	NaN	NaN	NaN	Na
EmployeeCount	1470.0	NaN	NaN	NaN	1.0	0.0	1.0	1.0	1.0	1.0	1
EmployeeNumber	1470.0	NaN	NaN	NaN	1024.865306	602.024335	1.0	491.25	1020.5	1555.75	2068
EnvironmentSatisfaction	1470.0	NaN	NaN	NaN	2.721769	1.093082	1.0	2.0	3.0	4.0	4
Gender	1470	2	Male	882	NaN	NaN	NaN	NaN	NaN	NaN	Na
HourlyRate	1470.0	NaN	NaN	NaN	65.891156	20.329428	30.0	48.0	66.0	83.75	100
JobInvolvement	1470.0	NaN	NaN	NaN	2.729932	0.711561	1.0	2.0	3.0	3.0	4
JobLevel	1470.0	NaN	NaN	NaN	2.063946	1.10694	1.0	1.0	2.0	3.0	5
JobRole	1470	9	Sales Executive	326	NaN	NaN	NaN	NaN	NaN	NaN	Na
JobSatisfaction	1470.0	NaN	NaN	NaN	2.728571	1.102846	1.0	2.0	3.0	4.0	4
MaritalStatus	1470	3	Married	673	NaN	NaN	NaN	NaN	NaN	NaN	Na
MonthlyIncome	1470.0	NaN	NaN	NaN	6502.931293	4707.956783	1009.0	2911.0	4919.0	8379.0	19999
MonthlyRate	1470.0	NaN	NaN	NaN	14313.103401	7117.786044	2094.0	8047.0	14235.5	20461.5	26999
NumCompaniesWorked	1470.0	NaN	NaN	NaN	2.693197	2.498009	0.0	1.0	2.0	4.0	9
Over18	1470	1	Y	1470	NaN	NaN	NaN	NaN	NaN	NaN	Na
OverTime	1470	2	No	1054	NaN	NaN	NaN	NaN	NaN	NaN	Na
PercentSalaryHike	1470.0	NaN	NaN	NaN	15.209524	3.659938	11.0	12.0	14.0	18.0	25
PerformanceRating	1470.0	NaN	NaN	NaN	3.153741	0.360824	3.0	3.0	3.0	3.0	4
RelationshipSatisfaction	1470.0	NaN	NaN	NaN	2.712245	1.081209	1.0	2.0	3.0	4.0	4
StandardHours	1470.0	NaN	NaN	NaN	80.0	0.0	80.0	80.0	80.0	80.0	80
StockOptionLevel	1470.0	NaN	NaN	NaN	0.793878	0.852077	0.0	0.0	1.0	1.0	3
TotalWorkingYears	1470.0	NaN	NaN	NaN	11.279592	7.780782	0.0	6.0	10.0	15.0	40
TrainingTimesLastYear	1470.0	NaN	NaN	NaN	2.79932	1.289271	0.0	2.0	3.0	3.0	6
WorkLifeBalance	1470.0	NaN	NaN	NaN	2.761224	0.706476	1.0	2.0	3.0	3.0	4
YearsAtCompany	1470.0	NaN	NaN	NaN	7.008163	6.126525	0.0	3.0	5.0	9.0	40
YearsInCurrentRole	1470.0	NaN	NaN	NaN	4.229252	3.623137	0.0	2.0	3.0	7.0	18
YearsSinceLastPromotion	1470.0	NaN	NaN	NaN	2.187755	3.22243	0.0	0.0	1.0	3.0	15

Numerical features like **Age**, **MonthlyIncome**, and **DistanceFromHome** have wide ranges. Categorical features show limited distinct categories suitable for encoding.

```
# -----
# Convert object columns to categorical type for efficiency
# -----
categorical_cols = df.select_dtypes(include='object').columns
for col in categorical_cols:
    df[col] = df[col].astype('category')

print("✅ Converted all object columns to categorical data types.")
print(df.dtypes)
```

✅ Converted all object columns to categorical data types.

Age	int64
Attrition	category
BusinessTravel	category
DailyRate	int64
Department	category

```

DistanceFromHome          int64
Education                  int64
EducationField             category
EmployeeCount              int64
EmployeeNumber              int64
EnvironmentSatisfaction    int64
Gender                     category
HourlyRate                 int64
JobInvolvement              int64
JobLevel                   int64
JobRole                     category
JobSatisfaction              int64
MaritalStatus               category
MonthlyIncome                int64
MonthlyRate                 int64
NumCompaniesWorked           int64
Over18                      category
OverTime                     category
PercentSalaryHike            int64
PerformanceRating            int64
RelationshipSatisfaction     int64
StandardHours                int64
StockOptionLevel              int64
TotalWorkingYears             int64
TrainingTimesLastYear         int64
WorkLifeBalance              int64
YearsAtCompany                int64
YearsInCurrentRole            int64
YearsSinceLastPromotion       int64
YearsWithCurrManager           int64
dtype: object

```

Converting to categorical helps with memory optimization and encoding operations later.

```

# -----
# Apply Label Encoding to all categorical features
# -----
le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col])

print("✓ Label Encoding Applied Successfully.")
df.head()

```

✓ Label Encoding Applied Successfully.

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	41	1	2	1102	2	1	2	1	1	1
1	49	0	1	279	1	8	1	1	1	1
2	37	1	2	1373	1	2	2	4	1	1
3	33	0	1	1392	1	3	4	1	1	1
4	27	0	2	591	1	2	1	3	1	1

5 rows × 35 columns

Label encoding converts categories into integers, making them usable in numeric models. Example: Gender → 0 for Female, 1 for Male.

```

# -----
# Standardize numeric columns for equal influence
# -----
num_cols = ['Age', 'MonthlyIncome', 'DistanceFromHome', 'YearsAtCompany', 'YearsInCurrentRole']
scaler = StandardScaler()
df[num_cols] = scaler.fit_transform(df[num_cols])

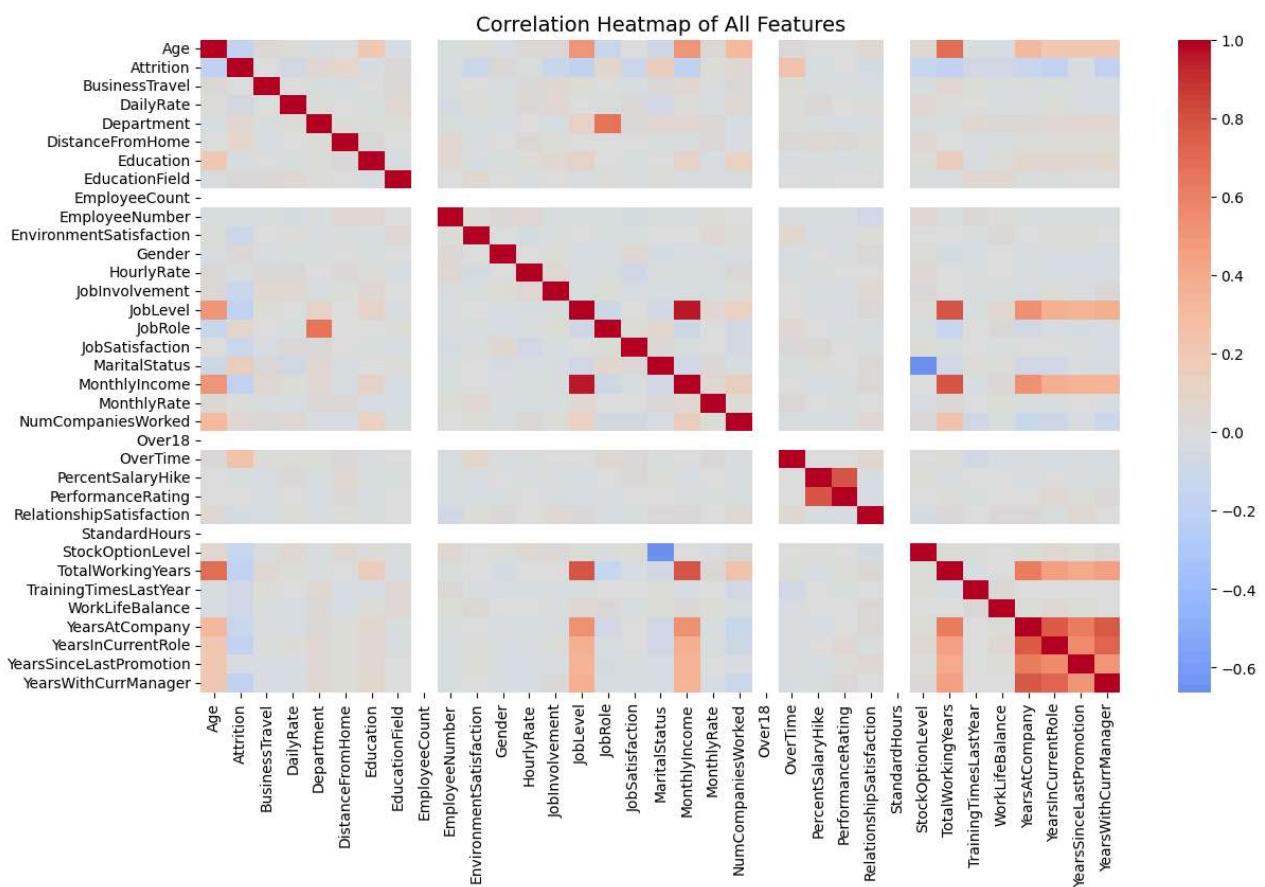
print("✓ Scaling completed for numerical features.")
df[num_cols].head()

```

Scaling completed for numerical features.

	Age	MonthlyIncome	DistanceFromHome	YearsAtCompany	YearsInCurrentRole
0	0.446350	-0.108350	-1.010909	-0.164613	-0.063296
1	1.322365	-0.291719	-0.147150	0.488508	0.764998
2	0.008343	-0.937654	-0.887515	-1.144294	-1.167687
3	-0.429664	-0.763634	-0.764121	0.161947	0.764998
4	-1.086676	-0.644858	-0.887515	-0.817734	-0.615492

```
# -----
# Plot correlation heatmap
# -----
plt.figure(figsize=(14,8))
sns.heatmap(df.corr(), cmap='coolwarm', center=0)
plt.title('Correlation Heatmap of All Features', fontsize=14)
plt.show()
```



From the heatmap, **JobLevel** and **MonthlyIncome** show high correlation. One of them can be dropped to reduce multicollinearity.

```
# -----
# Feature importance using Random Forest
# -----
X = df.drop('Attrition', axis=1)
y = df['Attrition']

model = RandomForestClassifier(random_state=42)
model.fit(X, y)

# Get top 10 important features
```

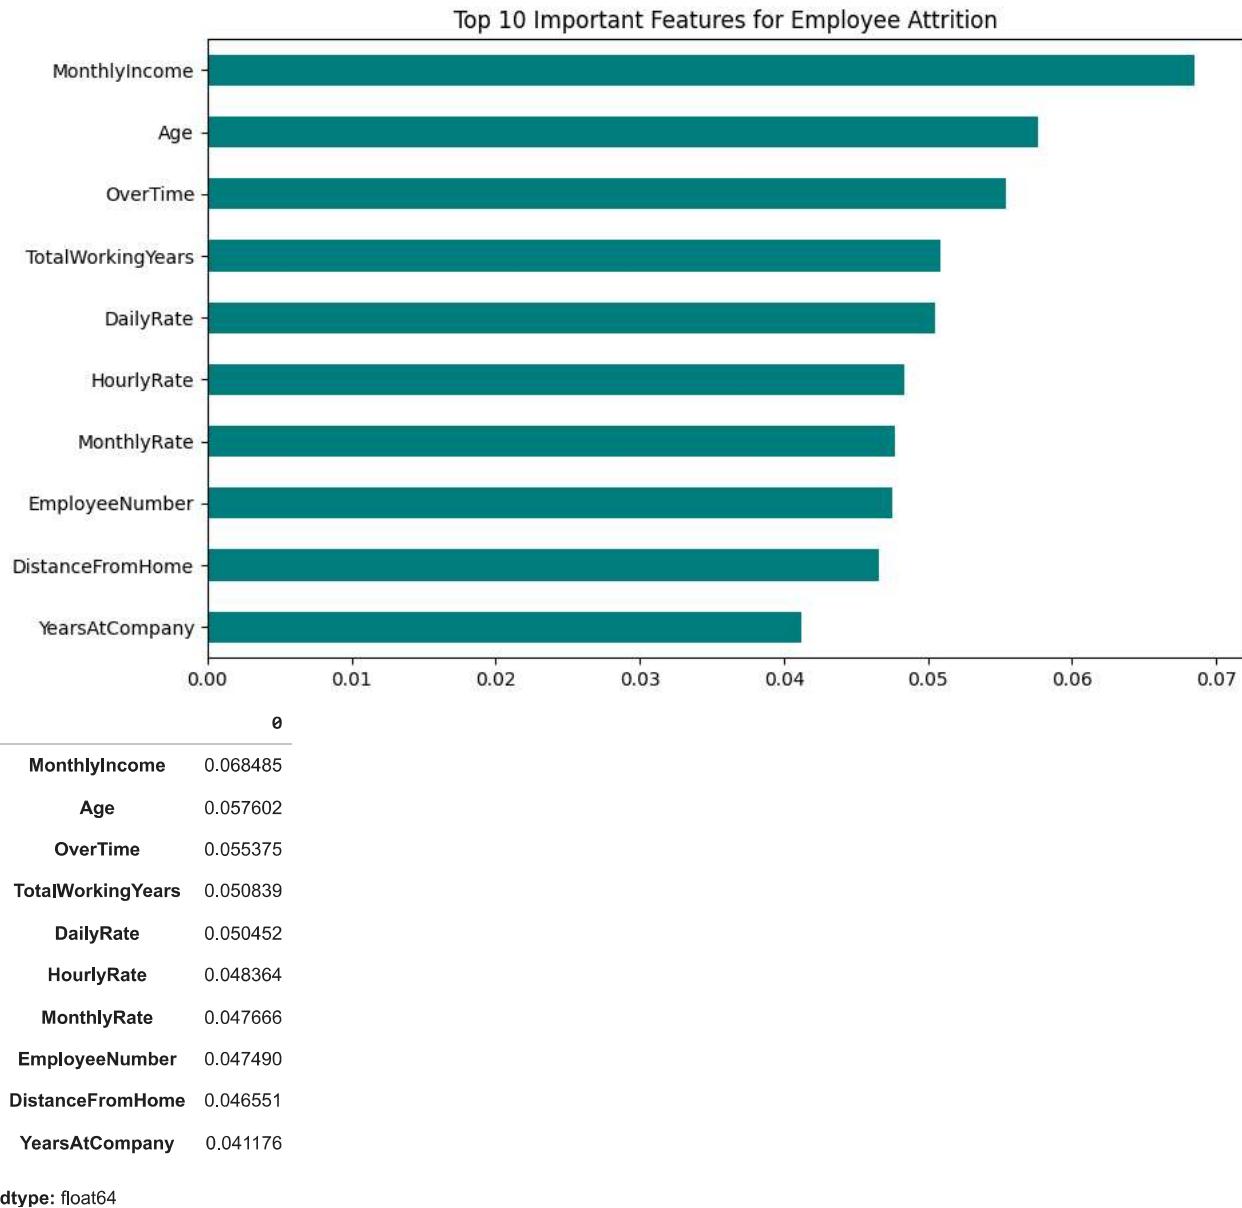
```

feat_imp = pd.Series(model.feature_importances_, index=X.columns).sort_values(ascending=False)

plt.figure(figsize=(10,6))
feat_imp.head(10).plot(kind='barh', color='teal')
plt.title('Top 10 Important Features for Employee Attrition')
plt.gca().invert_yaxis()
plt.show()

feat_imp.head(10)

```



The most important predictors include OverTime, MonthlyIncome, JobRole, and Age. *italicized text* These can be prioritized during model building or selection.

```

# -----
# Final summary of cleaned and transformed dataset
# -----
print("✓ Final Dataset Shape:", df.shape)
print("✓ No Missing Values:", df.isnull().sum().sum() == 0)
print("✓ Dataset Ready for Model Building")

df.head()

```

- Final Dataset Shape: (1470, 35)
- No Missing Values: True
- Dataset Ready for Model Building

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	En
0	0.446350	1	2	1102	2	-1.010909	2	1	1	1
1	1.322365	0	1	279	1	-0.147150	1	1	1	1
2	0.008343	1	2	1373	1	-0.887515	2	4	1	1
3	-0.429664	0	1	1392	1	-0.764121	4	1	1	1
4	-1.086676	0	2	591	1	-0.887515	1	3	1	1

5 rows × 35 columns

** Summary of Feature Engineering **

Dataset contained 1470 records and 35 columns.

No missing or duplicate values found.

Converted categorical data types for efficient memory use.

Applied Label Encoding and Standard Scaling to prepare features for modeling.